

# data1030-final

December 8, 2025

## 1 DATA1030 Final

Julien Song

### INSTRUCTIONS **EdStem:**

- Come up with a classification or regression problem!
- The goal of the final project is to develop a supervised ML pipeline from scratch by following the eight steps outlined in the lecture notes.
- There are three main reasons why a ML problem is difficult:
  - 1) missing data
  - 2) dataset is not IID (e.g., time series data, or one object is described by multiple data points - group structure)
  - 3) dataset is large (more than 100k points) so it is difficult to manage it on your laptop
- Choose a dataset with at least one of these difficulties!

**Rubric:** - Select a dataset and have an interesting/valuable/relevant classification or regression problem in mind to solve - The purpose of the semester-long project is to give you hands-on experience working with a machine learning pipeline. For the final presentation, the focus is on cross validation, the algorithms and parameters you tried, scores and uncertainties, and model explanations. The presentation is intended as a summary of your final project report.

- 1) EDA
- 2) splitting
- 3) preprocessing
- 4) evaluation metric
- 5) ML models
- 6) hyperparameter tuning
- 7) interpretability

**PROJECT Topic:** - predicting survival outcome of heart failure patients

**Data:** - PhysioNet - Hospitalized patients with heart failure: integrating electronic healthcare records and external outcome data - Zhang, Z., Cao, L., Zhao, Y., Xu, Z., Chen, R., Lv, L., & Xu, P. (2022). Hospitalized patients with heart failure: integrating electronic healthcare records and external outcome data (version 1.3). *PhysioNet*. RRID:SCR\_007345. <https://doi.org/10.13026/5m60-vs44> - **original publication:** Zhang Z, Cao L, Chen R, Zhao Y, Lv L, Xu Z, Xu P. Electronic health-care records and external outcome data for hospitalized patients with heart failure. *Sci Data*. 2021 Feb 5;8(1):46. doi: 10.1038/s41597-021-00835-9. PMID: 33547290; PMCID: PMC7865067 - patients who were admitted to a hospital in Sichuan, China between 2016 and 2019 - 168 variables for

2,008 patients with heart failure - Medications administered during the hospitalization are recorded in dat\_md.csv. - Subsequent admission & mortality = follow-up at 28 days, 3 months, 6 months  
-> classification problem (discrete ordinal)

```
[ ]: # IMPORT PACKAGES
import pandas as pd
import numpy as np

# EDA
import seaborn as sns
import matplotlib.pyplot as plt

# splitting
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold

# preprocessing
from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import LabelEncoder
from sklearn.utils.class_weight import compute_sample_weight
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
# from sklearn.ensemble import RandomForestRegressor
# from sklearn.ensemble import ExtraTreesRegressor
from sklearn.impute import KNNImputer

# models
from sklearn.pipeline import Pipeline
from imblearn.pipeline import Pipeline as ImbPipeline
from sklearn.model_selection import ParameterGrid
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
import xgboost

# metrics
from sklearn.metrics import f1_score

# interpretability
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import shap
```

```
[19]: # IMPORT DATA

# csv to df
df_og = pd.read_csv('heartfailure-ehr-data/dat.csv')

# display all
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)

# print(df_og.head())
# print(df_og.describe())
# print(df_og.shape) # (2008, 167)

[20]: # FINAL DF
df = df_og.copy()

# print(df.columns)

# create 1 y var:
# initial y options based on project goal of predicting survival outcome
drop_cols = ["Unnamed: 0", "inpatient.number", "outcome.during.hospitalization", "death.within.28.days", "re.admission.within.28.days", "death.within.3.months", "re.admission.within.3.months", "death.within.6.months", "re.admission.within.6.months", "time.of.death..days.from.admission.", "re.admission.time..days.from.admission.", "return.to.emergency.department.within.6.months", "time.to.emergency.department.within.6.months"]

# make y categorical -> death @ hosp, 28d, 3mo, 6mo, alive 6mo
y_series = pd.Series(name = 'survival.outcome')
for row in range(len(df)):
    if df.loc[row, "outcome.during.hospitalization"] == 'Dead':
        outcome = 'death.at.hospital'
    elif df.loc[row, "death.within.28.days"] == 1:
        outcome = 'death.within.28d'
    elif df.loc[row, "death.within.3.months"] == 1:
        outcome = 'death.within.3m'
    elif df.loc[row, "death.within.6.months"] == 1:
        outcome = 'death.within.6m'
    else:
        outcome = 'alive.after.6m'
    y_series[row] = outcome

# drop original y cols
df = df.drop(columns=drop_cols)

# drop column with all missing values:
```

```

# which columns completely missing:
# for col in df.columns:
#     if df[col].isnull().sum() == len(df):
#         print(col)
# cholinesterase

df = df.drop(columns=['cholinesterase'])

df['survival.outcome'] = y_series

```

[4]: # EDA  
# final dataset  
X = df.loc[:, df.columns != 'survival.outcome']  
y = df['survival.outcome']

# print(y.head())
print(y.shape) # (2008,)  
# # print(X.head())
print(X.shape) # (2008, 154)

(2008,)  
(2008, 153)

[21]: # EDA  
# check class balance:  
print(y.value\_counts(normalize=False))  
print(y.value\_counts(normalize=True))

# iid & imbalanced  
# -> kfold shuffle, class\_weight, f-score

```

survival.outcome
alive.after.6m      1951
death.within.28d    28
death.within.6m     15
death.at.hospital  11
death.within.3m     3
Name: count, dtype: int64

survival.outcome
alive.after.6m      0.971614
death.within.28d    0.013944
death.within.6m     0.007470
death.at.hospital  0.005478
death.within.3m     0.001494
Name: proportion, dtype: float64

```

[22]: # data types

```

# continuous:
continuous_ftrs = X.select_dtypes(include=['int64', 'float64']).columns.tolist()

# categorical:
all_cat_ftrs = X.select_dtypes(include=['object', 'category', 'bool']).columns
# ordinal vs. nominal:
# for col in all_cat_ftrs:
#     print(f"\n{col}: {X[col].unique()}")
# ordinal:
    # NYHA.cardiac.function.classification: ['III' 'II' 'IV']
    # Killip.grade: ['III' 'I' 'II' 'IV']
    # consciousness: ['Clear' 'ResponsiveToPain' 'ResponsiveToSound' ↴
    ↵ 'Nonresponsive']
    # respiratory.support.: [nan 'NIMV' 'IMV']
    # ageCat: ['(69,79]' '(59,69]' '(79,89]' '(49,59]' '(89,110]' '(29,39]' ↴
    ↵ '(39,49]' '(21,29]')
# binary (but want to assign 0/1):
    # admission.way: ['NonEmergency' 'Emergency']
    # type.II.respiratory.failure: ['NonTypeII' 'TypeII']
    # oxygen.inhalation: ['OxygenTherapy' 'AmbientAir']
ordinal_ftrs = ['NYHA.cardiac.function.classification', 'Killip.grade', ↴
    ↵ 'consciousness', 'respiratory.support.', 'ageCat',
    'admission.way', 'type.II.respiratory.failure', 'oxygen.
    ↵ inhalation']

ordinal_cats = [[['II', 'III', 'IV'], # NYHA.cardiac.function.classification
    ['I', 'II', 'III', 'IV'], # Killip.grade
    ['Clear', 'ResponsiveToSound', 'ResponsiveToPain', ↴
    ↵ 'Nonresponsive'], # consciousness
    ['NA', 'NIMV', 'IMV'], # respiratory.support.
    ['(21,29]', '(29,39]', '(39,49]', '(49,59]', '(59,69]', ↴
    ↵ '(69,79]', '(79,89]', '(89,110]'], # ageCat
    ['NonEmergency', 'Emergency'], # admission.way
    ['NonTypeII', 'TypeII'], # type.II.respiratory.failure
    ['AmbientAir', 'OxygenTherapy']] # oxygen.inhalation
]

# nominal:
nominal_ftrs = list(set(all_cat_ftrs) - set(ordinal_ftrs))

```

[7]: # EDA

```

# continuous = histogram
# ordinal = barplot (correct order)
# categorical = barplot
# continuous v. continuous = scatterplot, heatmap
# continuous v. categorical = category-specific histogram, boxplot, violin plot ↴
    (>3-4)

```

```
# categorical v. categorical = stacked barplot
```

[8]: # EDA  
# visualize individual variables:

```
# continuous = histogram  
# for col in continuous_ftrs:  
#     plt.figure(figsize=(4, 3)) # size to manage cell output  
#     plt.hist(X[col])#, kde=True)  
#     plt.title(f'Distribution of {col}')  
#     plt.xlabel(col)  
#     plt.ylabel('Frequency')  
#     plt.show()
```

[9]: # EDA  
# visualize individual variables:

```
# nominal = barplot  
# for col in nominal_ftrs:  
#     plt.figure(figsize=(4, 3)) # size to manage cell output  
#     sns.countplot(data=X, x=col)  
#     plt.title(f'Distribution of {col}')  
#     plt.xlabel(col)  
#     plt.ylabel('Count')  
#     plt.show()
```

[10]: # EDA  
# visualize individual variables:

```
# ordinal = barplot (in given order)  
# for col in ordinal_ftrs:  
#     plt.figure(figsize=(4, 3)) # size to manage cell output  
#     order = ordinal_cats[ordinal_ftrs.index(col)]  
#     sns.countplot(data=X, x=col, order=order)  
#     plt.title(f'Distribution of {col}')  
#     plt.xlabel(col)  
#     plt.ylabel('Count')  
#     plt.show()
```

[11]: # EDA  
# visualize outcome vs. features:

```
# categorical v. categorical = stacked bar plot  
# for col in nominal_ftrs:  
#     count_matrix = df.groupby([col, 'survival.outcome']).size().unstack()  
#     count_matrix_norm = count_matrix.div(count_matrix.sum(axis=1), axis=0)
```

```

#      count_matrix_norm.plot(kind='bar', stacked=True, cmap='Paired',□
↪figsize=(4, 3)) # size to manage cell output
#      plt.title(f'{col} vs. Survival Outcome')
#      plt.xlabel(col)
#      plt.ylabel('Count')
#      plt.legend(title='Survival Outcome', fontsize='small')
#      # plt.tight_layout()
#      plt.show()

# for col in ordinal_ftrs:
#     count_matrix = df.groupby([col, 'survival.outcome']).size().unstack()
#     count_matrix_norm = count_matrix.div(count_matrix.sum(axis=1), axis=0)

#     count_matrix_norm.plot(kind='bar', stacked=True, cmap='Paired',□
↪figsize=(4, 3)) # size to manage cell output
#     plt.title(f'{col} vs. Survival Outcome')
#     plt.xlabel(col)
#     plt.ylabel('Count')
#     plt.legend(title='Survival Outcome', fontsize='small')
#     # plt.tight_layout()
#     plt.show()

```

[12]:

```

# EDA
# continuous x v. outcome:

# categorical x continuous = category-specific histograms, box plot, violin plot
# for col in continuous_ftrs:
#     df[[col, 'survival.outcome']].boxplot(by='survival.outcome',□
↪figsize=(4,3)) # size to manage cell output
#     # plt.title(f'survival outcome by {col}')
#     plt.xlabel('survival outcome')
#     plt.ylabel(col)
#     plt.tight_layout()
#     plt.xticks(rotation=45)
#     plt.show()

```

[23]:

```

# EDA
# missing values
# how many missing values per column
missing_vals = df.isnull().sum()
missing_perc = (missing_vals / len(df)) * 100

# print(missing_vals)
# print(missing_perc)
# for col in df.columns:
#     miss = missing_vals[col]
#     perc = missing_perc[col]

```

```

#     if miss > 0:
#         print(f"{col}: {perc:.2f}%")

# cols with missing values 100-90, 80-90, 70-79, 60-69, 50-59, 40-49, 30-39,
# 20-29, 10-19, 1-9%
bins = [0, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100]
labels = ['0-9%', '10-19%', '20-29%', '30-39%', '40-49%', '50-59%', '60-69%', '70-79%', '80-89%', '90-100%']
df_missing = pd.DataFrame({'column': missing_vals.index, 'missing_count': missing_vals.values})
df_missing['missing_perc'] = (df_missing['missing_count'] / len(df_og)) * 100
df_missing['missing_bin'] = pd.cut(df_missing['missing_perc'], bins=bins, labels=labels, right=False)
missing_bin_counts = df_missing['missing_bin'].value_counts().sort_index()
print(missing_bin_counts)

# total missing values in dataset:
total_missing = missing_vals.sum()
total_values = df_og.size
overall_missing_perc = (total_missing / total_values) * 100
print(f"total missing values: {total_missing}")
print(f"overall missing percentage: {overall_missing_perc:.2f}%")

cont_missing = df[continuous_ftrs].isnull().sum().sum()
cont_missing_perc = (cont_missing / (len(continuous_ftrs) * len(df))) * 100
print(f"total missing values in continuous features: {cont_missing}")
print(f"total missing percentage in continuous features: {cont_missing_perc:.2f}%")
nom_missing = df[nominal_ftrs].isnull().sum().sum()
nom_missing_perc = (nom_missing / (len(nominal_ftrs) * len(df))) * 100
print(f"total missing values in nominal features: {nom_missing}")
print(f"total missing percentage in nominal features: {nom_missing_perc:.2f}%")
ord_missing = df[ordinal_ftrs].isnull().sum().sum()
ord_missing_perc = (ord_missing / (len(ordinal_ftrs) * len(df))) * 100
print(f"total missing values in ordinal features: {ord_missing}")
print(f"total missing percentage in ordinal features: {ord_missing_perc:.2f}%")

```

missing_bin	
0-9%	103
10-19%	7
20-29%	3
30-39%	1
40-49%	0
50-59%	25
60-69%	2
70-79%	4
80-89%	3

```
90-100%      6
Name: count, dtype: int64
total missing values: 57562
overall missing percentage: 17.17%
total missing values in continuous features: 55569
total missing percentage in continuous features: 19.91%
total missing values in nominal features: 27
total missing percentage in nominal features: 0.22%
total missing values in ordinal features: 1966
total missing percentage in ordinal features: 12.24%
```

```
[24]: # SPLITTING

# split during model iteration over multiple random states
random_states = [42, 52, 62, 72, 82]
```

```
[15]: # PREPROCESSING

# preprocess during model iteration over multiple random states
# random_states = [42, 52, 62, 72, 82]

# encode & impute
# continuous -> StandardScaler
# nominal -> OneHotEncoder
# ordinal -> OrdinalEncoder

# IterativeImputer not reasonable with iteration through such high # features
# -> use KNN Imputer
```

```
[28]: # MODELS

# logistic regression (classification)

# nominal transformer:
nominal_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                         ('onehot', OneHotEncoder(sparse_output=False, handle_unknown='ignore'))])
# ordinal transformer:
# We need to replace the NaN with a string first!
ordinal_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='NA')),
                                         ('ordinal', OrdinalEncoder(categories = ordinal_cats))])
# continuous transformer:
# KNN Imputer
numeric_transformer = Pipeline(steps=[('imputer', KNNImputer(n_neighbors=10)),
                                         ('scaler', StandardScaler())])
```

```

# preprocessor:
preprocessor = ColumnTransformer(transformers=[('num', numeric_transformer,
    ↪continuous_ftrs),
                                                ('cat', nominal_transformer,
    ↪nominal_ftrs),
                                                ('ord', ordinal_transformer,
    ↪ordinal_ftrs)])
```

# hyperparameter tuning:

```
lr_param_grid = {'logreg__penalty': ['l2'],
                 'logreg__C': np.logspace(-3, 0, 10),
                 'logreg__max_iter': [1000]}
```

# initialize result collection:

```
lr_test_scores = []
lr_models = []
```

# model:

```
for rs in random_states:
    # splitting: stratified kfold
    kf1 = StratifiedKFold(n_splits=3, shuffle=True, random_state=rs) #n 3
    ↪bc smallest class size = 3
    for other_idx, test_idx in kf1.split(X, y):
        X_other = X.iloc[other_idx]
        X_test = X.iloc[test_idx]
        y_other = y.iloc[other_idx] # y transformed
        y_test = y.iloc[test_idx]

    # pipeline:
    model = LogisticRegression(random_state = rs, class_weight='balanced')

    pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                               ('logreg', model)])
```

# cross validation:

```
kf2 = StratifiedKFold(n_splits=2, shuffle=True, random_state=rs) #n 2
↪bc smallest class size = 2
grid = GridSearchCV(pipeline, param_grid=lr_param_grid, scoring =
↪'f1_weighted', cv=kf2, return_train_score = True, verbose=True)
```

# model fitting:

```
grid.fit(X_other, y_other)
```

# results:

```
lr_models.append(grid)
```

```

y_test_pred = lr_models[-1].predict(X_test)
# test-score = f1_weighted
test_score = f1_score(y_test, y_test_pred, average='weighted')
lr_test_scores.append(test_score)

# best model:
lr_best_model_idx = np.argmax(lr_test_scores)
lr_best_model = lr_models[lr_best_model_idx]
lr_best_hyperparams = lr_best_model.best_params_
lr_best_test_score = lr_test_scores[lr_best_model_idx]

print("best model hyperparameters: ", lr_best_hyperparams)
print("best model test f1: ", lr_best_test_score)

# results (plot later)
lr_mean = np.mean(lr_test_scores)
lr_std = np.std(lr_test_scores)

```

Fitting 2 folds for each of 10 candidates, totalling 20 fits  
Fitting 2 folds for each of 10 candidates, totalling 20 fits  
Fitting 2 folds for each of 10 candidates, totalling 20 fits  
Fitting 2 folds for each of 10 candidates, totalling 20 fits  
Fitting 2 folds for each of 10 candidates, totalling 20 fits  
best model hyperparameters: {'logreg\_\_C': np.float64(1.0), 'logreg\_\_max\_iter': 1000, 'logreg\_\_penalty': 'l2'}  
best model test f1: 0.9447422523313181

```
[29]: # MODELS
# SVC

# same trasformers as log reg

# hyperparameter tuning:
svc_param_grid = {'svc__C': np.logspace(-3,0,10),
                  'svc__gamma': np.logspace(-3,0,10)}

# initialize result collection:
svc_test_scores = []
svc_models = []

# model:
for rs in random_states:
    # splitting: stratified kfold
    kf1 = StratifiedKFold(n_splits=3, shuffle=True, random_state=rs) #n 3 ↵bc smallest class size = 3
    for other_idx, test_idx in kf1.split(X, y):
        X_other = X.iloc[other_idx]
```

```

X_test = X.iloc[test_idx]
y_other = y.iloc[other_idx] # y transformed
y_test = y.iloc[test_idx]

# pipeline:
model = SVC(random_state = rs, class_weight='balanced')
pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                           ('svc', model)])

# cross validation:
kf2 = StratifiedKFold(n_splits=2, shuffle=True, random_state=rs) #n 2
↪bc smallest class size = 2
grid = GridSearchCV(pipeline, param_grid=svc_param_grid, scoring =
↪'f1_weighted',
                     cv=kf2, return_train_score = True, verbose=True)

# model fitting:
grid.fit(X_other, y_other)

# results:
svc_models.append(grid)
y_test_pred = svc_models[-1].predict(X_test)
# test-score = f1_weighted
test_score = f1_score(y_test, y_test_pred, average='weighted')
svc_test_scores.append(test_score)

# best model:
svc_best_model_idx = np.argmax(svc_test_scores)
svc_best_model = svc_models[svc_best_model_idx]
svc_best_hyperparams = svc_best_model.best_params_
svc_best_test_score = svc_test_scores[svc_best_model_idx]

print("best model hyperparameters: ", svc_best_hyperparams)
print("best model test f1: ", svc_best_test_score)

# results (plot later)
svc_mean = np.mean(svc_test_scores)
svc_std = np.std(svc_test_scores)

```

Fitting 2 folds for each of 100 candidates, totalling 200 fits  
Fitting 2 folds for each of 100 candidates, totalling 200 fits  
Fitting 2 folds for each of 100 candidates, totalling 200 fits  
Fitting 2 folds for each of 100 candidates, totalling 200 fits  
Fitting 2 folds for each of 100 candidates, totalling 200 fits  
best model hyperparameters: {'svc\_\_C': np.float64(0.46415888336127775),  
'svc\_\_gamma': np.float64(0.004641588833612777)}

```
best model test f1: 0.957290794077231
```

```
[30]: # MODELS
# RFC

# preprocessing: don't impute tree models -> trees only need preprocessing to
# numeric

# same categorical encoders as logreg & SVC
# # nominal encoder:
# nominal_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
#                                         ('onehot', OneHotEncoder(sparse_output=False,
#                                         handle_unknown='ignore'))])
# # ordinal encoder:
# ordinal_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='NA')),
#                                         ('ordinal', OrdinalEncoder(categories =
#                                         ordinal_cats))])
# continuous encoder: no imputation for tree models

# preprocessor:
preprocessor = ColumnTransformer(transformers=[('cat', nominal_transformer,
                                                nominal_ftrs),
                                                ('ord', ordinal_transformer,
                                                ordinal_ftrs)])

# hyperparameter tuning:
rfc_param_grid = {'rfc__n_estimators': [10,50,100,500],
                  'rfc__max_features': np.linspace(0.1,0.9,5),
                  'rfc__max_depth': [1,10,30]}

# initialize result collection:
rfc_test_scores = []
rfc_models = []

# model:
for rs in random_states:
    # splitting: kfold on other
    kf1 = StratifiedKFold(n_splits=3, shuffle=True, random_state=rs) #n 3
    #bc smallest class size = 3
    for other_idx, test_idx in kf1.split(X, y):
        X_other = X.iloc[other_idx]
        X_test = X.iloc[test_idx]
        y_other = y.iloc[other_idx] # y transformed
        y_test = y.iloc[test_idx]

    # pipeline:
```

```

    model = RandomForestClassifier(random_state=rs,
        ↪class_weight='balanced_subsample') # balanced_subsample for imbalanced
        pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                                    ('rfc', model)])

    # cross validation:
    kf2 = StratifiedKFold(n_splits=2, shuffle=True, random_state=rs) #n 2
    ↪bc smallest class size = 2
    grid = GridSearchCV(pipeline, param_grid=rfc_param_grid, scoring =
    ↪'f1_weighted',
                        cv=kf2, return_train_score = True, verbose=True)

    # model fitting:
    grid.fit(X_other, y_other)

    # results
    rfc_models.append(grid)
    y_test_pred = rfc_models[-1].predict(X_test)
    # test-score = f1_weighted
    test_score = f1_score(y_test, y_test_pred, average='weighted')
    rfc_test_scores.append(test_score)

# best model:
rfc_best_model_idx = np.argmax(rfc_test_scores)
rfc_best_model = rfc_models[rfc_best_model_idx]
rfc_best_hyperparams = rfc_best_model.best_params_
rfc_best_test_score = rfc_test_scores[rfc_best_model_idx]

print("best model hyperparameters: ", rfc_best_hyperparams)
print("best model test f1: ", rfc_best_test_score)

# results (plot later)
rfc_mean = np.mean(rfc_test_scores)
rfc_std = np.std(rfc_test_scores)

```

Fitting 2 folds for each of 60 candidates, totalling 120 fits  
Fitting 2 folds for each of 60 candidates, totalling 120 fits  
Fitting 2 folds for each of 60 candidates, totalling 120 fits  
Fitting 2 folds for each of 60 candidates, totalling 120 fits  
Fitting 2 folds for each of 60 candidates, totalling 120 fits  
best model hyperparameters: {'rfc\_\_max\_depth': 30, 'rfc\_\_max\_features':  
np.float64(0.9), 'rfc\_\_n\_estimators': 10}  
best model test f1: 0.9580779018335805

[31]: # MODELS  
# XGB

```

# preprocessing: don't impute tree models -> trees only need preprocessing to
# numeric
# same preprocessor as RFC
# # nominal encoder:
# nominal_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
#                                         ('onehot', OneHotEncoder(sparse_output=False, handle_unknown='ignore'))])
# # ordinal encoder:
# ordinal_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='NA')),
#                                         ('ordinal', OrdinalEncoder(categories =
#                                         ordinal_cats))])
# continuous encoder: no imputation for tree models

# # preprocessor:
# preprocessor = ColumnTransformer(transformers=[('cat', nominal_transformer,
#                                                 nominal_ftrs),
#                                                 ('ord', ordinal_transformer,
#                                                 ordinal_ftrs)])

# hyperparameter tuning:
xgb_param_grid = {'xgb__n_estimators': [10,50,100],
                  'xgb__max_depth': [5,30,50],
                  'xgb__learning_rate': [0.01, 0.05, 0.1],
                  'xgb__reg_lambda': np.logspace(-2,2,5),
                  'xgb__subsample': [0.3, 0.5, 0.7]}

# initialize result collection:
xgb_test_scores = []
xgb_models = []
xgb_test_sets = [] # best model -> save test set for cm

# y label encoder: (not in preprocessor)
# fit on the full set (y) to create a stable mapping label -> no leakage bc
# only records which string label maps to which integer, doesn't learn
# relationship between X & y
y_labenc = LabelEncoder()
y_labenc.fit(y)

# model:
for rs in random_states:
    # splitting: stratified kfold
    kf1 = StratifiedKFold(n_splits=3, shuffle=True, random_state=rs) #n 3
    #bc smallest class size = 3
    for other_idx, test_idx in kf1.split(X, y):

```

```

X_other = X.iloc[other_idx]
X_test = X.iloc[test_idx]
y_other = y.iloc[other_idx] # y transformed
y_test = y.iloc[test_idx]

# transform y other & test using label encoder fitted on full y
y_other_enc = y_labenc.transform(y_other)
y_test_enc = y_labenc.transform(y_test)

# pipeline:
# model = xgboost.XGBClassifier(early_stopping_rounds=40, enable_categorical=True) # true/false
model = xgboost.XGBClassifier(enable_categorical=True)
pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                           ('xgb', model)])

# cross validation:
kf2 = StratifiedKFold(n_splits=2, shuffle=True, random_state=rs) #n 2 bc smallest class size = 2
grid = GridSearchCV(pipeline, param_grid=xgb_param_grid, scoring = 'f1_weighted',
                     cv=kf2, return_train_score=True, verbose=True, error_score='raise')

# model fitting:
# no early stopping with GridSearchCV
class_weights = compute_sample_weight(class_weight='balanced', y=y_other_enc)
grid.fit(X_other, y_other_enc, xgb__sample_weight=class_weights, xgb__verbose=False) # verbose = validation_0-mlogloss output

# results
xgb_models.append(grid)
y_test_pred = xgb_models[-1].predict(X_test)

# test-score = f1_weighted
test_score = f1_score(y_test_enc, y_test_pred, average='weighted')
xgb_test_scores.append(test_score)
xgb_test_sets.append((X_test, y_test_enc))

# sanity check
# print(test_score)

# best model:
xgb_best_model_idx = np.argmax(xgb_test_scores)
xgb_best_model = xgb_models[xgb_best_model_idx]

```

```

xgb_best_hyperparams = xgb_best_model.best_params_
xgb_best_test_score = xgb_test_scores[xgb_best_model_idx]

print("best model hyperparameters: ", xgb_best_hyperparams)
print("best model test f1: ", xgb_best_test_score)

# results (plot later)
xgb_mean = np.mean(xgb_test_scores)
xgb_std = np.std(xgb_test_scores)

```

Fitting 2 folds for each of 405 candidates, totalling 810 fits  
Fitting 2 folds for each of 405 candidates, totalling 810 fits  
Fitting 2 folds for each of 405 candidates, totalling 810 fits  
Fitting 2 folds for each of 405 candidates, totalling 810 fits  
Fitting 2 folds for each of 405 candidates, totalling 810 fits  
best model hyperparameters: {'xgb\_\_learning\_rate': 0.05, 'xgb\_\_max\_depth': 30,  
'xgb\_\_n\_estimators': 50, 'xgb\_\_reg\_lambda': np.float64(1.0), 'xgb\_\_subsample':  
0.3}  
best model test f1: 0.9699978646166987

[47]: # BEST MODEL

```

# all 'best' models:
print("lr: ", lr_mean, " ", lr_std)
print("svc: ", svc_mean, " ", svc_std)
print("rfc: ", rfc_mean, " ", rfc_std)
print("xgb: ", xgb_mean, " ", xgb_std)

# xgb:
# best model hyperparameters: {'xgb__learning_rate': 0.05, 'xgb__max_depth': 30,  

# 'xgb__n_estimators': 50, 'xgb__reg_lambda': np.float64(0.01),  

# 'xgb__subsample': 0.5}
# best model test f1: 0.9600153034352713

best_model = xgb_best_model
best_X_test, best_y_test = xgb_test_sets[xgb_best_model_idx]

```

lr: 0.9411238583535871 , 0.003720000194119701  
svc: 0.9483896831638686 , 0.009935752032979972  
rfc: 0.9468757082403811 , 0.00931753233832749  
xgb: 0.9595519771332045 , 0.009438372944090704

[33]: # BASELINE

```

# multiclass classification: baseline = predict majority class for all points
# majority class = alive.after.6m = 0
majority_class = 0
y_baseline_pred = [majority_class] * len(best_y_test)

```

```
baseline_test_score = f1_score(best_y_test, y_baseline_pred, average='weighted')
print("baseline test f1: ", baseline_test_score)
```

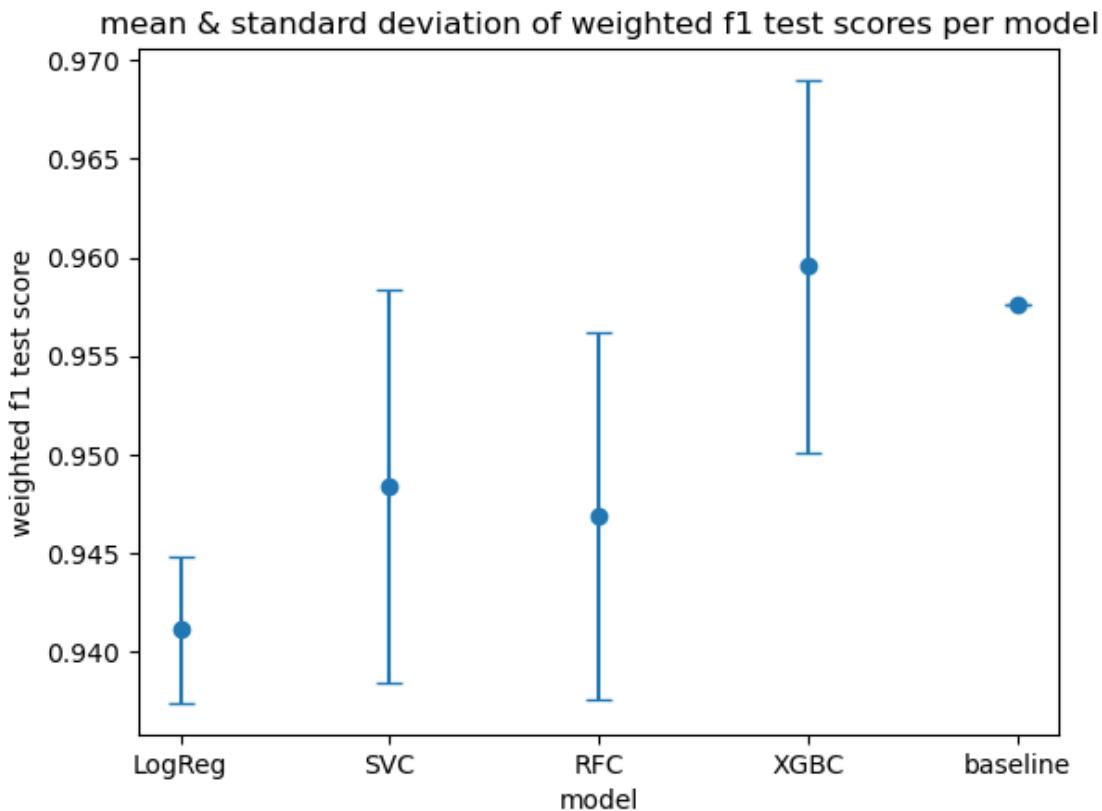
```
baseline test f1: 0.9576036563460791
```

```
[41]: # RESULTS
# plot means and standard deviations

labels = ['LogReg', 'SVC', 'RFC', 'XGBC', 'baseline']
means = [lr_mean, svc_mean, rfc_mean, xgb_mean, baseline_test_score]
stds = [lr_std, svc_std, rfc_std, xgb_std, 0] # assume baseline std=0

fig = plt.figure()
plt.errorbar(labels, means, yerr=stds, fmt='o', capsize=5, label='mean with std')
plt.xlabel('model')
plt.ylabel('weighted f1 test score')
plt.title('mean & standard deviation of weighted f1 test scores per model')
plt.show()

fig.savefig("mean_sd.png", dpi=300)
```

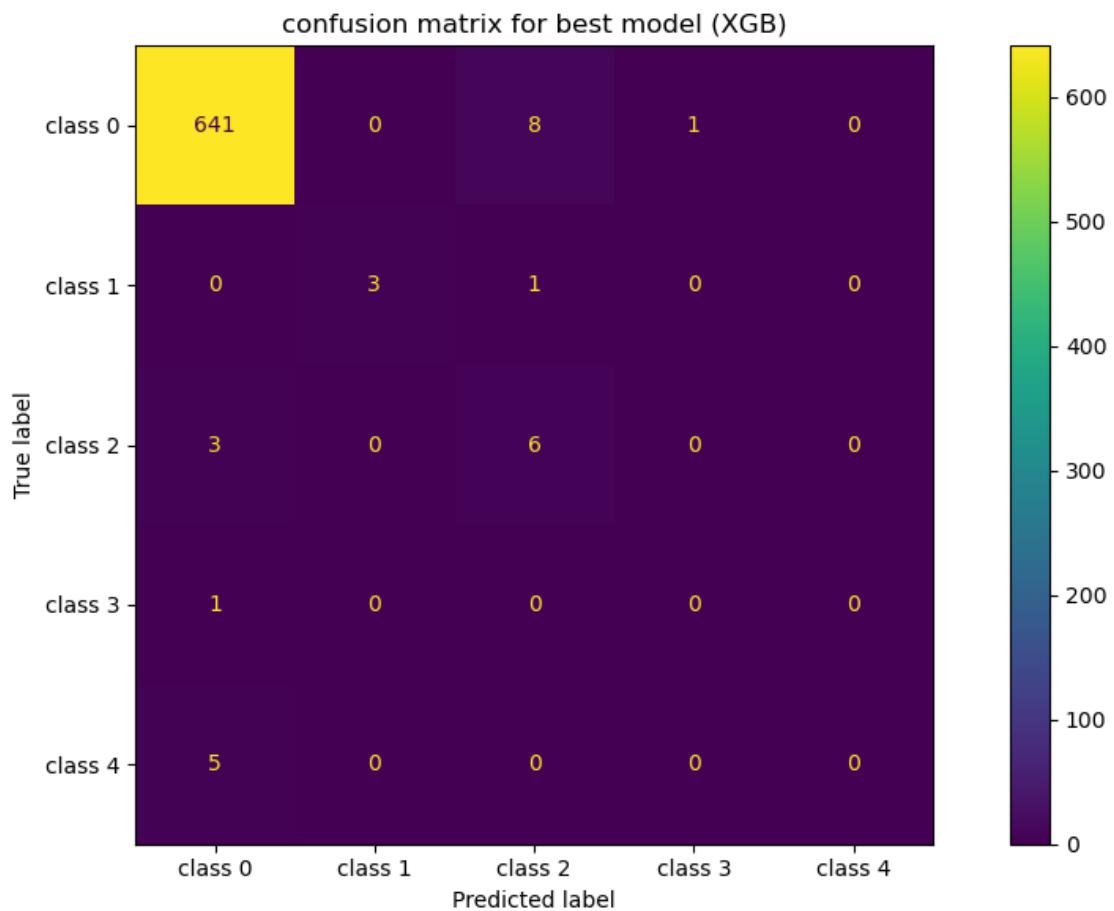


```
[39]: # INTERPRETABILITY
# confusion matrix
# best model:

# predict:
y_true = best_y_test
y_pred = best_model.predict(best_X_test)

# overall confusion matrix
cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(cm, display_labels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4'])
fig, ax = plt.subplots(figsize=(10,6))
disp.plot(ax=ax)
plt.title('confusion matrix for best model (XGB)')
plt.tight_layout()
plt.show()

fig.savefig("conf_matx.png", dpi=300)
```



```
[ ]: # INTERPRETABILITY
# global feature importances - xgb feature_importances_
best_model_step = best_model.best_estimator_.named_steps['xgb']
importances = best_model_step.feature_importances_
ftr_names = best_model.best_estimator_.named_steps['preprocessor'].
    ↪get_feature_names_out()
ftr_importance_df = pd.DataFrame({'feature': ftr_names,
                                    'importance': importances})
ftr_importance_df = ftr_importance_df.sort_values(by='importance', ↪
    ↪ascending=False) # sort descending

ftr_importance_df
```

	feature	importance
10	cat__DestinationDischarge_Died	0.108454
26	ord__respiratory.support.	0.084143
19	cat__discharge.department_Cardiology	0.066198
4	cat__occupation_missing	0.065827
12	cat__DestinationDischarge_Home	0.054816
6	cat__admission.ward_Cardiology	0.053379
1	cat__occupation_Others	0.050983
13	cat__DestinationDischarge_Unknown	0.046431
27	ord__ageCat	0.041249
28	ord__admission.way	0.039064
24	ord__Killip.grade	0.038888
14	cat__gender_Female	0.035087
7	cat__admission.ward_GeneralWard	0.032934
23	ord__NYHA.cardiac.function.classification	0.030319
3	cat__occupation_farmer	0.029714
20	cat__discharge.department_GeneralWard	0.027263
17	cat__type.of.heart.failure_Left	0.026680
29	ord__type.II.respiratory.failure	0.025906
11	cat__DestinationDischarge_HealthcareFacility	0.024827
21	cat__discharge.department_ICU	0.021101
25	ord__consciousness	0.020541
16	cat__type.of.heart.failure_Both	0.018168
2	cat__occupation_UrbanResident	0.017123
22	cat__discharge.department_Others	0.016507
8	cat__admission.ward_ICU	0.010975
9	cat__admission.ward_Others	0.007984
30	ord__oxygen.inhalation	0.005441
0	cat__occupation_Officer	0.000000
18	cat__type.of.heart.failure_Right	0.000000
5	cat__occupation_worker	0.000000
15	cat__gender_Male	0.000000

```
[120]: # INTERPRETABILITY
# local feature importances - SHAP
best_pipe = grid.best_estimator_

shap.initjs()
explainer = shap.TreeExplainer(best_pipe[1])
X_test_transformed = best_pipe[0].transform(best_X_test)
# shap_values = explainer.shap_values(X_test_transformed)

# print(np.shape(X_test_transformed)) # (669, 31)
# print(np.shape(shap_values)) # (669, 31, 5)

# index = 0 # the index of the point to explain
# print(explainer.expected_value[0]) # we explain class 0 predictions

feature_names = best_pipe[0].get_feature_names_out()
# shap.force_plot(explainer.expected_value[0], shap_values[index,:,:0], features_=X_test_transformed[index,:], feature_names = feature_names,
#                 matplotlib=True, show=False)
# # fig.title('SHAP force plot for class 0')
# # plt.show()
# # fig.savefig("shap.png", dpi=300)
# # Get the current figure and axes

# fig, ax = plt.gcf(), plt.gca()
# # fig.set_size_inches(12, 4)      # taller canvas

# for text in ax.texts:
#     # text.set_rotation(15) # Rotate by 45 degrees
#     # x, y = text.get_position()
#     # text.set_position((x, y - 0.5))    # push
#     if text.get_color() in ('#ff0052', '#008bff'):    # feature labels
#         text.set_rotation(45)
#     x, y = text.get_position()
#     text.set_position((x, y - 0.5))

# # plt.tight_layout()
# plt.show()

shap.summary_plot(shap_values[:, :, 1], X_test_transformed, feature_names = feature_names)
```

<IPython.core.display.HTML object>



```
[ ]: # local feature importances - SHAP
best_pipe = grid.best_estimator_

shap.initjs()
explainer = shap.TreeExplainer(best_pipe[1])
X_test_transformed = best_pipe[0].transform(best_X_test)
shap_values = explainer.shap_values(X_test_transformed)
```

```

# print(np.shape(X_test_transformed)) # (669, 31)
# print(np.shape(shap_values)) # (669, 31, 5)

print("expected value: ", explainer.expected_value[0])

for i in range(5):
    index = i # the index of the point to explain

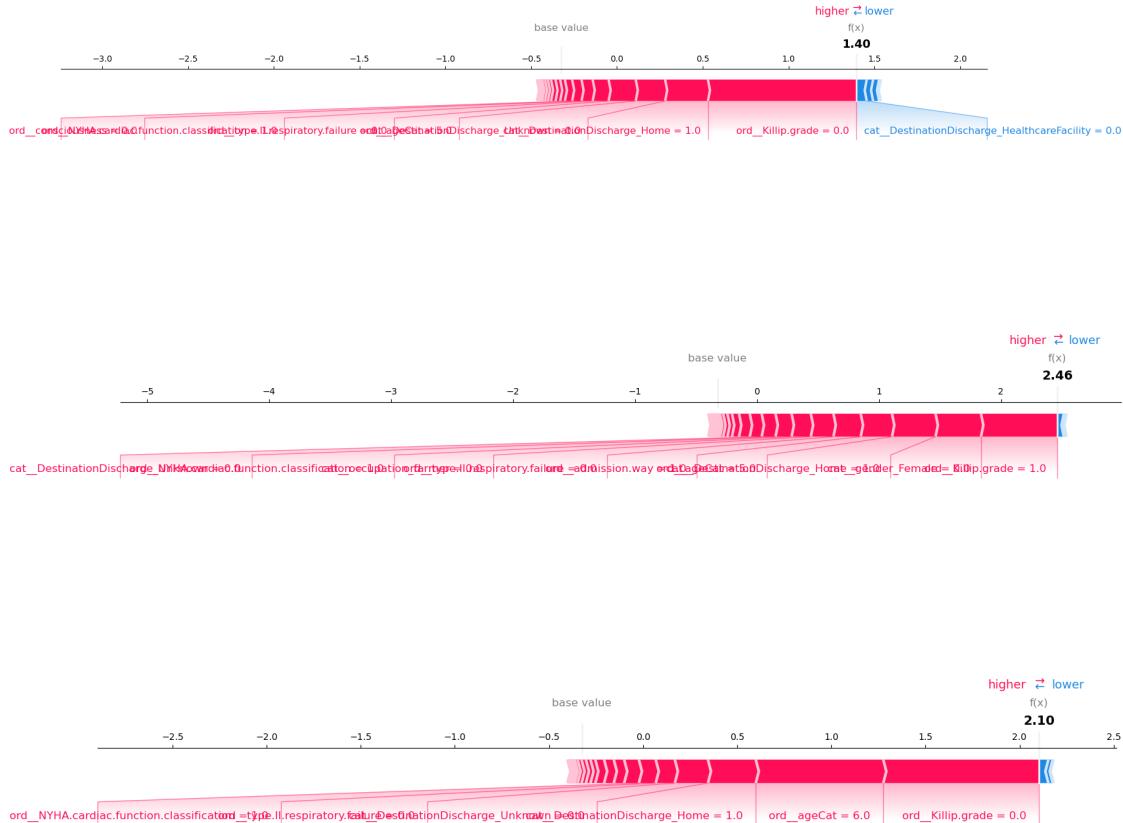
    feature_names = best_pipe[0].get_feature_names_out()
    shap.force_plot(explainer.expected_value[0], shap_values[index,:,:],  

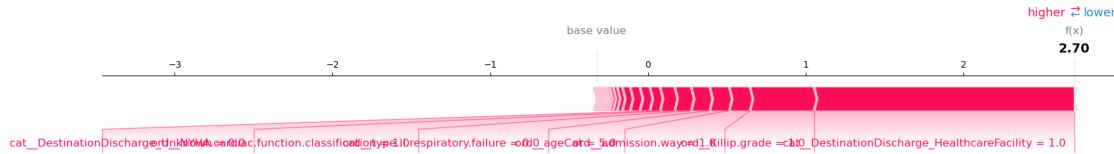
    ↪features = X_test_transformed[index,:], feature_names = feature_names,
    matplotlib=True)
    # fig.savefig(f"shap_{i}.png", dpi=300)

```

<IPython.core.display.HTML object>

expected value: -0.3249051





```
[140]: # pipeline: best_pipe[0] = preprocessor, best_pipe[1] = xgb estimator
best_est = best_pipe[1]
preproc = best_pipe[0]

# transform test set once
X_test_transformed = preproc.transform(best_X_test)      # numpy array
feature_names = preproc.get_feature_names_out()

# class names (if you used a LabelEncoder)

class_names = list(y_labenc.classes_)

# create explainer (TreeExplainer for tree models)
explainer = shap.TreeExplainer(best_est)
shap_values = explainer.shap_values(X_test_transformed)  # could be list or ndarray

# print(shap_values.shape)

# shap_values can be either:
# # - a list of arrays: shap_values[c] shape (n_samples, n_features)
# # - or an ndarray shape (n_samples, n_features, n_classes)
# if isinstance(shap_values, list):
# n_classes = len(shap_values)
#     # ensure each element is (n_samples, n_features)
# else:
#     # convert to list of per-class arrays for simpler indexing below
```

```

n_classes = shap_values.shape[2]
shap_values = [shap_values[:, :, c] for c in range(n_classes)]

sample_idx = 5 # the sample you want to explain

shap.initjs()
for c in range(n_classes):
    print(f"SHAP force plot for class {class_names[c]}:")
    base = explainer.expected_value[c]          # baseline for class c
    sv = shap_values[c][sample_idx]              # shape (n_features,)
    display(shap.force_plot(base, sv, X_test_transformed[sample_idx],
                            feature_names=feature_names))

# shap.initjs()
# for c in range(n_classes):
#     print(f"SHAP force plot for class {class_names[c]}:")
#     base = explainer.expected_value[c]          # baseline for class c
#     sv = shap_values[c][sample_idx]              # shape (n_features,)
#     display(shap.force_plot(base, sv, X_test_transformed[sample_idx],
#                             feature_names=feature_names, matplotlib=True,
#                             show=False))
#     # plt.xticks(rotation=45)
#     # plt.tight_layout()
#     # plt.show()

# for c in range(n_classes):
#     base = explainer.expected_value[c]
#     sv = shap_values[c][sample_idx]

#     # create the matplotlib force plot (returns a Matplotlib Figure)
#     fig = shap.force_plot(base, sv, X_test_transformed[sample_idx],
#                           feature_names=feature_names, matplotlib=True,
#                           show=False)

#     # rotate x tick labels on all axes in the figure
#     for ax in fig.axes:
#         for lbl in ax.get_xticklabels():
#             lbl.set_rotation(45)
#             lbl.set_ha('right')

#     plt.tight_layout()
#     display(fig)
#     plt.close(fig)

```

<IPython.core.display.HTML object>

SHAP force plot for class alive.after.6m:

```
<shap.plots._force.AdditiveForceVisualizer at 0x1386ab8f0>
SHAP force plot for class death.at.hospital:
<shap.plots._force.AdditiveForceVisualizer at 0x13a09d400>
SHAP force plot for class death.within.28d:
<shap.plots._force.AdditiveForceVisualizer at 0x13a09d400>
SHAP force plot for class death.within.3m:
<shap.plots._force.AdditiveForceVisualizer at 0x1383f5be0>
SHAP force plot for class death.within.6m:
<shap.plots._force.AdditiveForceVisualizer at 0x1384b6cc0>
```